


From observability to explainability with Datadog & Kubernetes MCP


Observability shows what happened

- **Logs** record events
- **Traces** follow a request
- **Monitoring** shows health and trends




Signals are not yet answers

- Human needs to connect the dot




LLMs add explainability

- Ask in plain language
- Collect relevant evidence
- Explain likely cause and next step



The Storedog demo

Goal: understand a running system and find faults sooner.



Demo 1: get the baseline

Show CPU and memory usage for Dogshop pods as a Markdown table. Add a simple comparison chart, highlight the three busiest pods, then inspect recent traffic-generator logs for errors. Do not make changes.

Demo 2: ads are unavailable

Run this first `make fault SCENARIO0=service-selector` and then send the prompt

Prompt

Why are Dogshop ads unavailable? Correlate the service, endpoints, pods, events, and Datadog signals. Show the evidence in a table, fix the cause, and verify recovery.

Demo 3: one cross-system summary

Prompt

Use Kubernetes MCP to summarize Dogshop workload health and Datadog MCP to find relevant monitors or telemetry. Return one incident-summary table with evidence, impact, and recommended action. Perform read-only operations only.

Other useful questions

- "Which workload has an image pull failure? Explain the warning events and repair it."
- "Why is the frontend rollout not completing? Inspect the failed readiness probe."
- "Scale ads to three replicas and verify the rollout."
- "Delete one ads pod and verify that the Deployment recovers."

Closing remarks

- Observability gives us evidence
- MCP gives an LLM controlled access to that evidence
- Explainability helps teams understand and respond sooner